

Name: Leonardo Neves

Andrew ID: lneves

Machine Learning for Text Mining

Homework 3

1. Statement of Assurance

I assure that I am responsible for all the code and answers for this homework.

2. Experiments

- a) Describe the custom weighting scheme that you have implemented. Explain your motivation for creating this weighting scheme.

My custom weighting method started as the weighted harmonic mean as I thought it would have a similar result as the weighted sum, but giving a higher effect to lower scores. Thus, I realized that the page ranks and the relevance scores were not on the same order of magnitude and, to better weight them, I have taken the log of their absolute value. I got the minus because the relevance scores are negative.

$$CM = - \frac{W1 + W2}{\frac{w1}{\log(abs(PR))} + \frac{w2}{\log(abs(PageRelevanceScore))}}$$

- b) Report of the performance of the 9 approaches.

I. Metric: MAP

Method \ Weighting Scheme	NS	WS	CM
GPR	0.0458	0.2642	0.2637
QTSPR	0.0456	0.2644	0.2637
PTSPR	0.0469	0.2644	0.2638

II. Metric: Precision at 11 standard recall levels

Method \ Weighting Scheme	NS	WS	CM
GPR	ircl_prn.0.00 all 0.1447	ircl_prn.0.00 all 0.8417	ircl_prn.0.00 all 0.8409
	ircl_prn.0.10 all 0.0874	ircl_prn.0.10 all 0.5996	ircl_prn.0.10 all 0.5929
	ircl_prn.0.20 all 0.0787	ircl_prn.0.20 all 0.4757	ircl_prn.0.20 all 0.4737
	ircl_prn.0.30 all 0.0738	ircl_prn.0.30 all 0.3765	ircl_prn.0.30 all 0.3789

	ircl_prn.0.40 all 0.0701 ircl_prn.0.50 all 0.0655 ircl_prn.0.60 all 0.0535 ircl_prn.0.70 all 0.0301 ircl_prn.0.80 all 0.0115 ircl_prn.0.90 all 0.0074 ircl_prn.1.00 all 0.0041 P5 all 0.0368 P10 all 0.0447 P15 all 0.0474 P20 all 0.0474 P30 all 0.0421 P100 all 0.0450 P200 all 0.0518 P500 all 0.0601 P1000 all 0.0301	ircl_prn.0.40 all 0.3134 ircl_prn.0.50 all 0.2411 ircl_prn.0.60 all 0.1688 ircl_prn.0.70 all 0.0921 ircl_prn.0.80 all 0.0541 ircl_prn.0.90 all 0.0390 ircl_prn.1.00 all 0.0098 P5 all 0.6000 P10 all 0.5105 P15 all 0.4509 P20 all 0.4145 P30 all 0.3596 P100 all 0.1939 P200 all 0.1275 P500 all 0.0601 P1000 all 0.0301	ircl_prn.0.40 all 0.3143 ircl_prn.0.50 all 0.2427 ircl_prn.0.60 all 0.1675 ircl_prn.0.70 all 0.0916 ircl_prn.0.80 all 0.0567 ircl_prn.0.90 all 0.0390 ircl_prn.1.00 all 0.0103 P5 all 0.6000 P10 all 0.5132 P15 all 0.4456 P20 all 0.4132 P30 all 0.3579 P100 all 0.1942 P200 all 0.1274 P500 all 0.0601 P1000 all 0.0301
QTSPR	ircl_prn.0.00 all 0.1444 ircl_prn.0.10 all 0.0867 ircl_prn.0.20 all 0.0779 ircl_prn.0.30 all 0.0733 ircl_prn.0.40 all 0.0695 ircl_prn.0.50 all 0.0650 ircl_prn.0.60 all 0.0533 ircl_prn.0.70 all 0.0302 ircl_prn.0.80 all 0.0114 ircl_prn.0.90 all 0.0073 ircl_prn.1.00 all 0.0040 P5 all 0.0368 P10 all 0.0447 P15 all 0.0491 P20 all 0.0474 P30 all 0.0421 P100 all 0.0453 P200 all 0.0517 P500 all 0.0601 P1000 all 0.0301	ircl_prn.0.00 all 0.8417 ircl_prn.0.10 all 0.5994 ircl_prn.0.20 all 0.4756 ircl_prn.0.30 all 0.3773 ircl_prn.0.40 all 0.3148 ircl_prn.0.50 all 0.2417 ircl_prn.0.60 all 0.1689 ircl_prn.0.70 all 0.0920 ircl_prn.0.80 all 0.0542 ircl_prn.0.90 all 0.0391 ircl_prn.1.00 all 0.0098 P5 all 0.6000 P10 all 0.5105 P15 all 0.4509 P20 all 0.4145 P30 all 0.3596 P100 all 0.1939 P200 all 0.1274 P500 all 0.0601 P1000 all 0.0301	ircl_prn.0.00 all 0.8402 ircl_prn.0.10 all 0.5928 ircl_prn.0.20 all 0.4735 ircl_prn.0.30 all 0.3785 ircl_prn.0.40 all 0.3146 ircl_prn.0.50 all 0.2428 ircl_prn.0.60 all 0.1674 ircl_prn.0.70 all 0.0916 ircl_prn.0.80 all 0.0567 ircl_prn.0.90 all 0.0388 ircl_prn.1.00 all 0.0103 P5 all 0.6000 P10 all 0.5132 P15 all 0.4456 P20 all 0.4145 P30 all 0.3579 P100 all 0.1937 P200 all 0.1278 P500 all 0.0601 P1000 all 0.0301
PTSPR	rcl_prn.0.00 all 0.1451 ircl_prn.0.10 all 0.0832	ircl_prn.0.00 all 0.8417 ircl_prn.0.10 all 0.6004	ircl_prn.0.00 all 0.8403 ircl_prn.0.10 all 0.5935

	ircl_prn.0.20 all 0.0785	ircl_prn.0.20 all 0.4760	ircl_prn.0.20 all 0.4739
	ircl_prn.0.30 all 0.0744	ircl_prn.0.30 all 0.3774	ircl_prn.0.30 all 0.3787
	ircl_prn.0.40 all 0.0700	ircl_prn.0.40 all 0.3144	ircl_prn.0.40 all 0.3145
	ircl_prn.0.50 all 0.0617	ircl_prn.0.50 all 0.2419	ircl_prn.0.50 all 0.2425
	ircl_prn.0.60 all 0.0491	ircl_prn.0.60 all 0.1684	ircl_prn.0.60 all 0.1675
	ircl_prn.0.70 all 0.0270	ircl_prn.0.70 all 0.0921	ircl_prn.0.70 all 0.0916
	ircl_prn.0.80 all 0.0114	ircl_prn.0.80 all 0.0543	ircl_prn.0.80 all 0.0567
	ircl_prn.0.90 all 0.0072	ircl_prn.0.90 all 0.0390	ircl_prn.0.90 all 0.0388
	ircl_prn.1.00 all 0.0041	ircl_prn.1.00 all 0.0099	ircl_prn.1.00 all 0.0104
	P5 all 0.0263	P5 all 0.6000	P5 all 0.6000
	P10 all 0.0421	P10 all 0.5132	P10 all 0.5132
	P15 all 0.0491	P15 all 0.4526	P15 all 0.4456
	P20 all 0.0474	P20 all 0.4145	P20 all 0.4145
	P30 all 0.0456	P30 all 0.3605	P30 all 0.3579
	P100 all 0.0497	P100 all 0.1939	P100 all 0.1942
	P200 all 0.0580	P200 all 0.1276	P200 all 0.1275
	P500 all 0.0601	P500 all 0.0601	P500 all 0.0601
	P1000 all 0.0301	P1000 all 0.0301	P1000 all 0.0301

III. Metric: Wall-clock running time in seconds

Method \ Weighting Scheme	NS	WS	CM
GPR	0.2117162315	0.1956176820	0.1972206078
QTSPR	0.2232007290	0.2210894509	0.2128124237
PTSPR	0.2002020195	0.2124132419	0.2261456753

IV. Parameters

For the GPR, I have used PRW of 1000 and relevance of 0.5 for the WS and PRW of -8 and relevanceW of 100 for the CM.

For the QTSPR, I have used beta as 0.02 and gamma as 0.18 and the same values for WS and PRW.

For PTSPR, I have used beta = 1.0 and gamma = 0 for the NS and beta = 0.04 and gamma as 0.16 for WS and CM. Same values for the weights of WS and PRW for the GPR were used.

- c) Compare these 9 approaches based on the various metrics described above.

All experiments were based on PR after 10 iterations. We can see that the No Search approaches have a worse result. This makes clear that the PageRank alone is not a really good metric for ranking documents. On the other hand, the search relevance alone has a baseline of 0.2636 MAP and all the results that combine the PR and the SearchRelevance data are better than that, showing that PR is able to improve performance when combined.

The query-based page rank did not have a much better result. Without the search, it shows a worse result as we could expect, since there is no real query being searched, while it got a better result for the Weighted Sum approach. The best results were the ones that used the personalized TSPR, getting higher precision and recall using the weighted sum. The Custom Method, a adaptation of the weighted harmonic mean, got a better precision on higher ranks than the weighted sum on most experiments, maybe making it more suitable for web search engines where the top ranks are more important, despite of having a MAP.

Because of the well-structured design and by caching the offline PageRank computation, we did not see much difference among computation time of any of the algorithms, all of them being fast enough to be used by a user.

- d) Analyze these various algorithms, parameters, and discuss your general observations about using PageRank algorithms.

As we could see on question C, PageRank is not a good metric alone or, at least, is not better than using the search relevance alone, but it might be used to improve the results. Personalized Topic Sensitive PageRank improved the results significantly and, even without the search relevance it showed better than the other two. For the parameters, since the PR value is much lower than the search relevance value, we had to tune the weights in order to get a better result and that could get them to be in the same order of magnitude. Also, using PTSPR asked for higher parameters for beta, showing that the teleportation matrix is more important for this configuration than for QTSPR.

- e) 1. What could be some novel ways for search engines to estimate whether a query can benefit from personalization?

We could consider the query to be a really small document. Also, we could pre-process documents of the same topic, maybe using the kmeans algorithm, to get a good vector to represent the given topic. Then, we could compute the similarity between the query and this representation and, if some vectors showed a high similarity to the query, it might indicate that personalization would be a good idea.

2. What could be some novel ways of identifying the user's interests (e.g. the user's topical interest distribution $\Pr(t|u)$) in general?

One obvious but not so effective way would be to have a set of labeled documents and explicitly ask user for feedback. This has the pro of being 100% sure about the user interest but most users would not like to have those questions often. A better approach might be to monitor the user queries. If his searches frequently retrieve documents labeled as a topic, we might think he is interested on that particular topic. Also, if he searches for documents, selects one and stays a reasonable time reading that document, it might represent interest as well. This time spent and the amount of documents of a particular topic normally retrieved by the user might be good for estimating initial probabilities that could be tuned by treating this problem as a classification problem. We could then, from the assumptions, use our knowledge to rank the documents differently and observe the user response to that. If we retrieve a document based on our assumptions and the user do spend a reasonable time reading that document, we can assume it was a positive prediction and assume it was bad otherwise.

3. Details of the software implementation

- a) Describe your design decisions and high-level software architecture;

I have created a function to use the Power method for computing the offline matrix of pagerank. My transition matrix was created using a sparse matrix from Scipy. I have used a map of maps to map every user to every query and every user-query pair to the topics distribution. I have also used a map to store the pagerank matrices based on the topic.

- b) Describe major data structures and any other data structures you used for speeding up the computation of PageRank;

Page rank computation was fast because I have used Sparse matrix to create the M matrix and and the normalize function from sklearn to normalize the rows by the out edges values. Sparse matrix multiplication is very fast using those libraries.

- c) Describe any programming tools or libraries that you used;

I have used Scipy for sparse matrix computations, Numpy for handling data and sklearn for normalizing the rows.

- d) Describe strengths and weaknesses of your design, and any problems that your system encountered

I think my design allows me to make computation very fast. It is easy to see on the experiments that the computation is so fast that it is hard to determine which method is faster. I have also cached the offline page rank computation, which made my code faster for avoiding unnecessary computation.

4. Describe how to run your code (programming environment, command line, etc.)

To run the code, run *pagerank.py method path-to-distro alpha beta-factor pagerankWeight searchRelevanceWeight*

Method = “WS”, “CM” or “NS”

Path-to-distro = “query-topic-distro.txt” or “user-topic-distro.txt”

Alpha = dampening factor = 0.8

Beta-factor = proportion of the remaining value $(1 - \alpha) = 0.2$ that should go to beta. The rest will go to gamma.

PageRankWeight and searchRelevanceWeight are based for the WS or CM methods.

Code assumes the indri-lists folder is in the same folder as the code.